Assignment 3

Due Date: Mentioned in eLearning

Instructions

- You should write code on Databricks for the two parts below. Clearly demarcate between the two parts.

- You need to have a cover sheet, which should include the names and netids of every team member.

- You are allowed to work in pairs i.e. a group of two students is allowed. Please write the names of the group members on the cover page.

- **You have a total of 4 free late days for the entire semester. You can use at most 2 days for any one assignment. After four days have been used up, there will be a penalty of 10% for each late day. The submission for this assignment will be closed 2 days after the due date.**

- Please ask all questions on Piazza, not via email.

1 PageRank for Airports

PageRank algorithm can be used to evaluate relative importance of nodes in a connected environment. It is based on the concept of in-links and out-links of a node and is used to rank nodes in a graph in order of their importance. Details about this algorithm and its implementation using MapReduce can be found in Chapter 5 of the reference book
Data Intensive Text Processing using MapReduce.
You can also look at the slides available at
http://lintool.github.io/UMD-courses/bigdata-2015-Spring/slides/session05.pdf.

Note that you cannot use any external library that automatically computes PageRank, including Spark GraphX.

The dataset for this project will be a graph that shows connections between various airports. This data is available at:Bureau of Transportation website. You would need to do the following to download the data:

1. Go to https://transtats.bts.gov/

2. On the left menu, click under "Aviation" under the "By Mode" block.

3. On the next page, click Air Carrier Statistics (Form 41 Traffic)- U.S. Carriers

4. On the next page, click download link under T-100 Domestic Segment (U.S. Carriers)

5. On the next page, set Filter Year = Most Recent Year, Filter Period = Any month on which data is available (e.g. July), and select following fields:

   - Origin (Origin Airport Code)
   - OriginCityName (optional)
   - Dst (Destination Airport Code)
   - DstCityName (optional)

6. Download and unzip to get a csv file.

Below are the requirements of the project:

1. Your program should read the input data from a location on Databricks.

2. You will compute the page rank of each node (airport) based on number of inlinks and outlinks. There may be multiple connections between two airports, and you should consider them independent of each other to compute the number of inlinks and outlinks. For example, if node A is connected to node B with an out-count of 10 and node C with an out-count of 10, then the total number of outlinks for node A would be 20.

3. You have to limit yourself to maximum number of 20 iterations.

4. You will use the following equation to compute PageRank:

$$\text{PR}(x) = \alpha \times \frac{1}{N} + (1 - \alpha) \times \sum_{1}^{n} \frac{PR(t_i)}{C(t_i)}$$

where $\alpha = 0.15$ and x is a page with inlinks from $t_1, t_2, \ldots, t_n$, $C(t)$ is the out-degree of t, and N is the total number of nodes in the graph.


Figure 1: PageRank computation

You can initialize all the PageRank values to be 10.0.

5. You have to display the top 10 airports with the highest PageRank on the screen. If every PageRank is the same, it means you didn't run the code properly.

2 Tweet Processing & Classification using Pipelines

In this part, we will work with a set of Tweets about US airlines and examine their sentiment polarity. More details about the dataset is available at:
https://www.kaggle.com/crowdflower/twitter-airline-sentiment. It is part of the Kaggle competition on Twitter US Airline Sentiment. Our aim is to learn to classify Tweets as either "positive", "neutral", or "negative" by using logistic regression classifier and pipelines for pre-processing and model building.

You will do this using Databricks notebook, with the input file being hosted on Databricks also.

You need to create a pipeline with the following steps. Again, you need to create a pipeline and not have to run these steps individually. Below are the steps of the project:

1. **Loading**: First step is to load the text file from the path specified in argument 1. After that, you will need to remove rows where the *text* field is *null*.

2. **Pre-Processing**: You will start by creating a pre-processing step with the following stages:

   - **Stop Word Remover**: Remove stop-words from the text column
     Hint: Use the org.apache.spark.ml.feature.StopWordsRemover class.

   - **Tokenizer**: Transform the *text* column into words by breaking down the sentence into words .
     Hint: Use the import org.apache.spark.ml.feature.Tokenizer class.

   - **Term Hashing**: Convert words to term-frequency vectors
     Hint: Use the import org.apache.spark.ml.feature.HashingTF class

   - **Label Conversion**: The label is a string e.g. "Positive", which you need to convert to numeric format
     Hint: Use the import org.apache.spark.ml.feature.StringIndexer class

   Remember that you need to create a pipeline of the above steps and then transform the raw input dataset to a pre-processed dataset.

3. **Model Creation** - You will need to create a logistic regression classification model. You will have to create a **ParameterGridBuilder** for parameter tuning and then use the **CrossValidator** object for finding the best model parameters. More details can be seen here:
   https://spark.apache.org/docs/2.2.0/api/scala/index.html

4. **Model Testing & Cross Validation**: Next, you will need to train and test your model on the given dataset and output classification evaluation metrics, such as accuracy, etc. You can see details of multi-class evaluation metrics at
   https://spark.apache.org/docs/2.2.0/mllib-evaluation-metrics.html.

5. **Output**: Finally, you have to write the output the classification metrics to the screen.